

CS 238P | OPERATING SYSTEMS | Homework 2

[Name: Shefali Gupta Student ID: 57806943]

Exercise 1:

[GENERAL OBSERVATIONS]

To find start of kernel address, from one terminal run make qemu-nox

```
— vagrant@andromeda-68:/vagrant/cs238p/xv6-public — ssh shefalg@andromeda-68.ics.uci.edu
[vagrant@andromeda-68:/vagrant/cs238p/xv6-public$ make qemu-nox
dd if=/dev/zero of=xv6.img count=10000
10000+0 records in
10000+0 records out
5120000 bytes (5.1 MB) copied, 1.27406 s, 4.0 MB/s
dd if=bootblock of=xv6.img conv=notrunc
1+0 records in
1+0 records out
512 bytes (512 B) copied, 0.00103448 s, 495 kB/s
dd if=kernel of=xv6.img seek=1 conv=notrunc
327+1 records in
327+1 records out
167488 bytes (167 kB) copied, 0.0852635 s, 2.0 MB/s
qemu-system-i386 -nographic -drive file=fs.img,index=1,media=disk,format=raw -drive file=xv6.img,index=0,media=disk,format=raw -smp 2 -m 512
xv6...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ ]
```

From another terminal, run: nm kernel | grep _start

```
[vagrant@andromeda-68:/vagrant/cs238p/xv6-public$ nm kernel | grep _start
8010a48c D _binary_entryother_start
8010a460 D _binary_initcode_start
0010000c T _start
vagrant@andromeda-68:/vagrant/cs238p/xv6-public$
```

We found the address of kernel start as 0x10000c.

Inside vagrant, doing make qemu-nox-gdb.

```

~ — vagrant@andromeda-68: /vagrant/cs238p/xv6-public — ssh shefali@andromeda-68.ics.uci.edu
vagrant@andromeda-68: /vagrant/cs238p/xv6-public$ make qemu-nox-gdb
dd if=/dev/zero of=xv6.img count=10000
10000+0 records in
10000+0 records out
5120000 bytes (5.1 MB) copied, 1.19758 s, 4.3 MB/s
dd if=bootblock of=xv6.img conv=notrunc
1+0 records in
1+0 records out
512 bytes (512 B) copied, 0.001261 s, 406 kB/s
dd if=kernel of=xv6.img seek=1 conv=notrunc
327+1 records in
327+1 records out
167488 bytes (167 kB) copied, 0.0018023 s, 2.0 MB/s
*** Now run 'gdb'.
qemu-system-i386 -nographic -drive file=fs.img,index=1,media=disk,format=raw -drive file=xv6.img,index=0,media=disk,format=raw -smp 2 -m 512 -S -gdb tcp::260
01

```

Switching to another terminal and running gdb. We put a breakpoint on start address (0x10000c). Checking the registers.

```

Terminal Shell Edit View Window Help
shefali — vagrant@andromeda-68: /vagrant/cs238p/xv6-public
GNU gdb (Ubuntu/Linaro 7.4-2012.04-0ubuntu2.1) 7.4-2012.04
Copyright (C) 2012 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type "show copying"
and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
For bug reporting instructions, please see:
<http://bugs.launchpad.net/gdb-linaro/>.
+ target remote localhost:26001
warning: A handler for the OS ABI "GNU/Linux" is not built into this configuration
of GDB. Attempting to continue with the default i8086 settings.

The target architecture is assumed to be i8086
[f000:ffff] 0xfffff0: ljmp $0xf000,$0xe05b
0x0000ffff0 in ?? ()
+ symbol-file kernel
(gdb) br * 0010000c
Invalid number "0010000c".
(gdb) br * 0x10000c
Breakpoint 1 at 0x10000c
(gdb) info reg
eax          0x0      0
ecx          0x0      0
edx          0x663    1635
ebx          0x0      0
esp          0x0      0x0
ebp          0x0      0x0
esi          0x0      0
edi          0x0      0
eip          0xfffff0  0xfffff0
eflags        0x2      [ ]
cs           0xf000   61440
ss           0x0      0
ds           0x0      0
es           0x0      0
fs           0x0      0
gs           0x0      0
(gdb)

```

We see that, only the edx, cd (code segment) and eip (instruction pointer) registers point to an address, to fetch and execute the very first instruction of xv6 booting, inside bootblock.asm

Also, the stack is initially empty.

```
(gdb) x/24x $esp
0x0:    0x00000000      0x00000000      0x00000000      0x00000000
0x10:   0x00000000      0x00000000      0x00000000      0x00000000
0x20:   0x00000000      0x00000000      0x00000000      0x00000000
0x30:   0x00000000      0x00000000      0x00000000      0x00000000
0x40:   0x00000000      0x00000000      0x00000000      0x00000000
0x50:   0x00000000      0x00000000      0x00000000      0x00000000
(gdb)
```

Right before call to bootmain, the stack pointer register (esp) is initialized.

movl \$start, %esp ; %esp= 0x7c00

Now putting a breakpoint on 0x7c00 (where call to bootmain() occurs) , and inspecting the state of registers for every instruction executed until we hit the breakpoint.

```
0x0000ffff in ?? ()
+ symbol-file kernel
(gdb) br * 0x7c00
Breakpoint 1 at 0x7c00
(gdb) info reg
eax      0x0      0
ecx      0x0      0
edx      0x663    1635
ebx      0x0      0
esp      0x0      0x0
ebp      0x0      0x0
esi      0x0      0
edi      0x0      0
eip      0xffff0  0xffff0
eflags     0x2      [ ]
cs       0xf000   61440
ss       0x0      0
ds       0x0      0
es       0x0      0
fs       0x0      0
gs       0x0      0
(gdb) si
[f000:e05b] 0xfe05b: cmpl    $0x0,%cs:0x6ac8
0x0000e05b in ?? ()
(gdb) si
[f000:e062] 0xfe062: jne     0xfd2e1
0x0000e062 in ?? ()
(gdb) si
[f000:e066] 0xfe066: xor     %dx,%dx
0x0000e066 in ?? ()
(gdb) info reg
eax      0x0      0
ecx      0x0      0
edx      0x663    1635
ebx      0x0      0
esp      0x0      0x0
ebp      0x0      0x0
esi      0x0      0
edi      0x0      0
eip      0xe066   0xe066
eflags     0x46      [ PF ZF ]
cs       0xf000   61440
ss       0x0      0
ds       0x0      0
es       0x0      0
fs       0x0      0
gs       0x0      0
(gdb)
```

```

gs          0x0      0
(gdb) si
[f000:e05b] 0xfe05b: cmpl    $0x0,%cs:0x6ac8
0x0000e05b in ?? ()
(gdb) si
[f000:e062] 0xfe062: jne     0xfd2e1
0x0000e062 in ?? ()
(gdb) si
[f000:e066] 0xfe066: xor     %dx,%dx
0x0000e066 in ?? ()
(gdb) info reg
eax          0x0      0
ecx          0x0      0
edx          0x663    1635
ebx          0x0      0
esp          0x0      0x0
ebp          0x0      0x0
esi          0x0      0
edi          0x0      0
eip          0xe066   0xe066
eflags        0x46    [ PF ZF ]
cs           0xf000   61440
ss           0x0      0
ds           0x0      0
es           0x0      0
fs           0x0      0
gs           0x0      0
(gdb) si
[f000:e068] 0xfe068: mov     %dx,%ss
0x0000e068 in ?? ()
(gdb) info reg
eax          0x0      0
ecx          0x0      0
edx          0x0      0
ebx          0x0      0
esp          0x0      0x0
ebp          0x0      0x0
esi          0x0      0
edi          0x0      0
eip          0xe068   0xe068
eflags        0x46    [ PF ZF ]
cs           0xf000   61440
ss           0x0      0
ds           0x0      0
es           0x0      0
fs           0x0      0
gs           0x0      0
(gdb) ■

```

Before calling `mov $0x7000, %esp`, state of registers is as:

```

(gdb) si
[f000:e06a] 0xfe06a: mov     $0x7000,%esp
0x0000e06a in ?? ()
(gdb) info reg
eax          0x0      0
ecx          0x0      0
edx          0x0      0
ebx          0x0      0
esp          0x0      0x0
ebp          0x0      0x0
esi          0x0      0
edi          0x0      0
eip          0xe06a   0xe06a
eflags        0x46    [ PF ZF ]
cs           0xf000   61440
ss           0x0      0
ds           0x0      0
es           0x0      0
fs           0x0      0
gs           0x0      0
(gdb) ■

```

After the instruction is executed, the Stack register (esp) points to an address location of 0x7000.

```
0x00000000 11: . . .
(gdb) info reg
eax          0x0      0
ecx          0x0      0
edx          0x0      0
ebx          0x0      0
esp          0x7000  0x7000
ebp          0x0      0x0
esi          0x0      0
edi          0x0      0
eip          0xe070  0xe070
eflags        0x46    [ PF ZF ]
cs           0xf000  61440
ss           0x0      0
ds           0x0      0
es           0x0      0
fs           0x0      0
gs           0x0      0
(gdb) $
```

Similary, before call to 0xf0574, state of registers and stack is as follows:

```
0x00000000 11: . . .
=> 0xf09ab:   call   0xf0574
0x000f09ab in ?? ()
(gdb) info reg
eax          0xf5b58  1006424
ecx          0x6fd4   28628
edx          0xf447b  1000571
ebx          0x0      0
esp          0x6fcc   0x6fcc
ebp          0x0      0x0
esi          0x0      0
edi          0x0      0
eip          0xf09ab  0xf09ab
eflags        0x2      [ ]
cs           0x8      8
ss           0x10     16
ds           0x10     16
es           0x10     16
fs           0x10     16
gs           0x10     16
(gdb) x/24x $esp
0x6fcc: 0x000f34da  0x000f447b  0x000f5b5c  0x00000000
0x6fdc: 0x00000000  0x00000000  0x00000000  0x00000000
0x6fec: 0x00000000  0x00000000  0x00000000  0x00000000
0x6ffc: 0x00000000  0x00000000  0x00000000  0x00000000
0x700c: 0x00000000  0x00000000  0x00000000  0x00000000
0x701c: 0x00000000  0x00000000  0x00000000  0x00000000
(gdb) $
```

Now, proceeding so on, right before bootmain is called, at 0x7c00 (our breakpoint), state of registers and stack is:

```
Breakpoint 1, 0x00007c00 in ?? ()  
|(gdb) info reg  
eax          0xaa55  43605  
ecx          0x0      0  
edx          0x80    128  
ebx          0x0      0  
esp          0x6f20  0x6f20  
ebp          0x0      0x0  
esi          0x0      0  
edi          0x0      0  
eip          0x7c00  0x7c00  
eflags       0x202   [ IF ]  
cs           0x0      0  
ss           0x0      0  
ds           0x0      0  
es           0x0      0  
fs           0x0      0  
gs           0x0      0  
|(gdb) x/24x $esp  
0x6f20: 0xf000d239  0x00000000  0x00006f62  0x00000000  
0x6f30: 0x00008d0d  0x00008cd4  0x00000000  0x00000000  
0x6f40: 0x00006ae5  0x00007c00  0x00007c00  0x00000080  
0x6f50: 0x0000f1c5d 0x000f3c9a  0x00000000  0x00007c00  
0x6f60: 0x00000000  0x00000000  0x00000000  0x00000000  
0x6f70: 0x00000000  0x00800000  0x00000000  0xaa550000  
|(gdb) █
```

EIP register holds value of next instruction to be executed, which is 0x7c00. ESP points to 0x6f20.

Before calling bootmain at address 0x7d34, ESP points to 0x7c00

```
^--> 0x7c48:     call   0x7d34  
0x00007c48 in ?? ()  
|(gdb) info reg  
eax          0x0      0  
ecx          0x0      0  
edx          0x80    128  
ebx          0x0      0  
esp          0x7c00  0x7c00  
ebp          0x0      0x0  
esi          0x0      0  
edi          0x0      0  
eip          0x7c48  0x7c48  
eflags       0x6      [ PF ]  
cs           0x8      8  
ss           0x10    16  
ds           0x10    16  
es           0x10    16  
fs           0x0      0  
gs           0x0      0  
|(gdb) █
```

[ANSWERS TO QUESTIONS]

Question: Where in bootasm.S is the stack pointer initialized?

Answer:

In bootasm.S, right before call to bootmain, the stack pointer register (esp) is initialized.

```
    movl $start, %esp ; // %esp= 0x7c00
```

```
# Set up the stack pointer and call into C.  
movl    $start, %esp  
call    bootmain
```

Question: Single step through the call to bootmain; what is on the stack now?

Answer: When the IP points to the bootmain subroutine, the content in the stack is now **0x7c4d**, which is the next instruction address.

(State of registers and top of stack before call to bootmain at 0x7d34)

```
=> 0x7c48:      call   0x7d34  
0x000007c48 in ?? ()  
|(gdb) info reg  
eax          0x0      0  
ecx          0x0      0  
edx          0x80     128  
ebx          0x0      0  
esp          0x7c00   0x7c00  
ebp          0x0      0x0  
esi          0x0      0  
edi          0x0      0  
eip          0x7c48   0x7c48  
eflags        0x6      [ PF ]  
cs           0x8      8  
ss           0x10     16  
ds           0x10     16  
es           0x10     16  
fs           0x0      0  
gs           0x0      0  
|(gdb) x/24x $esp  
0x7c00: 0x8ec031fa  0x8ec08ed8  0xa864e4d0  0xb0fa7502  
0x7c10: 0xe464e6d1  0x7502a864  0xe6dfb0fa  0x16010f60  
0x7c20: 0x200f7c78  0xc88366c0  0xc0220f01  0x087c31ea  
0x7c30: 0x10b86600  0x8ed88e00  0x66d08ec0  0x8e0000b8  
0x7c40: 0xbce88ee0  0x00007c00  0x0000e7e8  0x00b86600  
0x7c50: 0xc289668a  0xb866ef66  0xef668ae0  0x9066feeb  
(gdb) █
```

(State of registers and top of stack, right after call to bootmain)

```
|(gdb) si
=> 0x7d34:      push    %ebp
0x00007d34 in ?? ()
|(gdb) info reg
eax          0x0      0
ecx          0x0      0
edx          0x80     128
ebx          0x0      0
esp          0x7bfc   0x7bfc
ebp          0x0      0x0
esi          0x0      0
edi          0x0      0
eip          0x7d34   0x7d34
eflags        0x6      [ PF ]
cs           0x8      8
ss           0x10     16
ds           0x10     16
es           0x10     16
fs           0x0      0
gs           0x0      0
|(gdb) x/24x $esp
0x7bfc: 0x00007c4d 0x8ec031fa 0x8ec08ed8 0xa864e4d0
0x7c0c: 0xb0fa7502 0xe464e6d1 0x7502a864 0xe6dfb0fa
0x7c1c: 0x16010f60 0x200f7c78 0xc88366c0 0xc0220f01
0x7c2c: 0x087c31ea 0x10b86600 0x8ed88e00 0x66d08ec0
0x7c3c: 0x8e0000b8 0xbce88ee0 0x00007c00 0x0000e7e8
0x7c4c: 0x00b86600 0xc289668a 0xb866ef66 0xef668ae0
|(gdb)
```

Question: What do first the assembly instructions of bootmain do to the stack?

Answer: push %ebp → saves the content of %ebp (0) to the top of the stack.

```
=> 0x7d34:      push    %ebp
0x00007d34 in ?? ()
|(gdb) info reg
eax          0x0      0
ecx          0x0      0
edx          0x80     128
ebx          0x0      0
esp          0x7bfc   0x7bfc
ebp          0x0      0x0
esi          0x0      0
edi          0x0      0
eip          0x7d34   0x7d34
eflags        0x6      [ PF ]
cs           0x8      8
ss           0x10     16
ds           0x10     16
es           0x10     16
fs           0x0      0
gs           0x0      0
|(gdb) x/24x $esp
0x7bfc: 0x00007c4d 0x8ec031fa 0x8ec08ed8 0xa864e4d0
0x7c0c: 0xb0fa7502 0xe464e6d1 0x7502a864 0xe6dfb0fa
0x7c1c: 0x16010f60 0x200f7c78 0xc88366c0 0xc0220f01
0x7c2c: 0x087c31ea 0x10b86600 0x8ed88e00 0x66d08ec0
0x7c3c: 0x8e0000b8 0xbce88ee0 0x00007c00 0x0000e7e8
0x7c4c: 0x00b86600 0xc289668a 0xb866ef66 0xef668ae0
|(gdb) si
=> 0x7d35:      mov     %esp,%ebp
0x00007d35 in ?? ()
|(gdb) x/24x $esp
0x7bf8: 0x00000000 0x00007c4d 0x8ec031fa 0x8ec08ed8
0x7c08: 0xa864e4d0 0xb0fa7502 0xe464e6d1 0x7502a864
0x7c18: 0x6dfb0fa 0x16010f60 0x200f7c78 0xc88366c0
0x7c28: 0xc0220f01 0x087c31ea 0x10b86600 0x8ed88e00
0x7c38: 0x66d08ec0 0x8e0000b8 0xbce88ee0 0x00007c00
0x7c48: 0x0000e7e8 0x00b86600 0xc289668a 0xb866ef66
|(gdb)
```

Question: Look in bootmain in bootblock.asm for the call that changes eip to 0x10000c. What does the call do to the stack?

Answer: From bootblock.asm, we find that kernel entry() function is called at **0x7dbc**. The compiled assembly instructions show that entry address of kernel was saved at the **address 0x10018**. We put a breakpoint at 0x7dbc and execute all instructions, until we hit this breakpoint. State of registers and stack before the call is:

```
(gdb) br * 0x7dbc
Breakpoint 2 at 0x7dbc
(gdb) c
Continuing.
=> 0x7dbc:    call    *0x10018

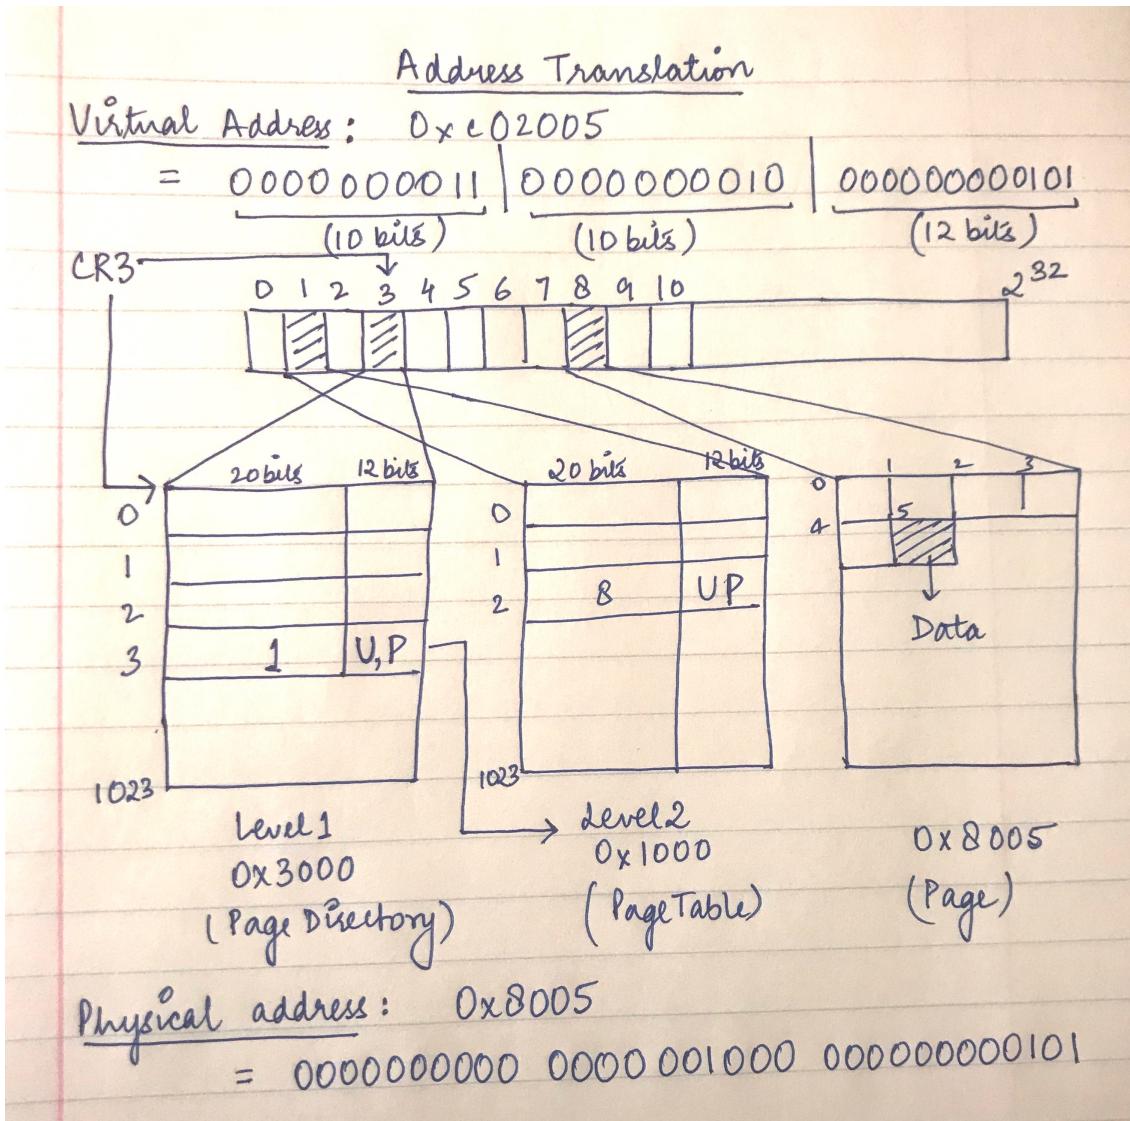
Breakpoint 2, 0x000007dbc in ?? ()
(gdb) info reg
eax      0x0      0
ecx      0x0      0
edx      0x1f0    496
ebx      0x10074  65652
esp      0x7bc0   0x7bc0
ebp      0x7bf8   0x7bf8
esi      0x0      0
edi      0x1154a8 1135784
eip      0x7dbc   0x7dbc
eflags   0x46    [ PF ZF ]
cs       0x8      8
ss       0x10     16
ds       0x10     16
es       0x10     16
fs       0x0      0
gs       0x0      0
(gdb) x/24x $esp
0x7bc0: 0x00000000 0x00000000 0x00000000 0x00000000
0x7bd0: 0x00000000 0x00000000 0x00000000 0x00010074
0x7be0: 0x00000000 0x00000000 0x00000000 0x00000000
0x7bf0: 0x00000000 0x00000000 0x00000000 0x00007c4d
0x7c00: 0x8ec031fa 0x8ec08ed8 0xa864e4d0 0xb0fa7502
0x7c10: 0xe464e6d1 0x7502a864 0xe6dfb0fa 0x16010f60
(gdb) █
```

After call,

```
(gdb) si
=> 0x1000c:    mov    %cr4,%eax
0x001000c in ?? ()
(gdb) info reg
eax      0x0      0
ecx      0x0      0
edx      0x1f0    496
ebx      0x10074  65652
esp      0x7bbc   0x7bbc
ebp      0x7bf8   0x7bf8
esi      0x0      0
edi      0x1154a8 1135784
eip      0x10000c 0x10000c
eflags   0x46    [ PF ZF ]
cs       0x8      8
ss       0x10     16
ds       0x10     16
es       0x10     16
fs       0x0      0
gs       0x0      0
(gdb) x/24x $esp
0x7bbc: 0x000007dc2 0x00000000 0x00000000 0x00000000
0x7bcc: 0x00000000 0x00000000 0x00000000 0x00000000
0x7bcd: 0x00010074 0x00000000 0x00000000 0x00000000
0x7bec: 0x00000000 0x00000000 0x00000000 0x00000000
0x7bfc: 0x000007c4d 0x8ec031fa 0x8ec08ed8 0xa864e4d0
0x7c0c: 0xb0fa7502 0xe464e6d1 0x7502a864 0xe6dfb0fa
(gdb) █
```

Here, we see that the call to 0x10018 jumps to kernel code address 0x10000c. The return address of C function call entry() is also saved on stack (0x000007dc2).

Exercise 2:



During the translation of virtual address to physical address, the virtual address of 32 bits is split into 10, 10 and 12 bits. The first 10 bits indicate the index in level 1 page table (Page directory), which in this example is 3 in the physical memory address 0x3000. Since, level 2 page table table address is given as (0x1000), this means entry at 3 in PDE is 1. The second chunk of 10 bits in the virtual address, give the index in the level 2 page table, where address of final page table is present. In this case, it is 2. Now, since we are given final page in physical memory as having address 0x8005, this means entry at index 2 in level 2 page table, points to page at index 8 in the physical memory. The last 12 bits of the virtual address(5 in this case), indicate the index in the physical page, where the data is present.

Exercise 3:

Before kvmalloc() is called:

```
(qemu) info pg
VPN range      Entry      Flags      Physical page
[00000-003ff]  PDE[000]   --S-A---WP 00000-003ff
[80000-803ff]  PDE[200]   --S-A---WP 00000-003ff
(qemu)
```

We set a breakpoint in main and hit next, until kvmalloc() is called.

```
(gdb) b main
Breakpoint 1 at 0x80102e70: file main.c, line 19.
(gdb) c
Continuing.
The target architecture is assumed to be i386
=> 0x80102e70 <main>: push %ebp

Breakpoint 1, main () at main.c:19
19
(gdb) n
=> 0x80102e7a <main+10>:      movl $0x80400000,0x4(%esp)
20      kinit1(end, P2V(4*1024*1024)); // phys page allocator
(gdb) n
=> 0x80102e8e <main+30>:      call 0x80106b30 <kvmalloc>
21      kvmalloc(); // kernel page table
(gdb) n
=> 0x80102e93 <main+35>:      call 0x80103000 <mpinit>
22      mpinit(); // detect other processors
(gdb)
```

Snapshot of Page table after kvmalloc() is called:

```
(qemu) info pg
VPN range      Entry      Flags      Physical page
[80000-803ff]  PDE[200]   ----A--UWP
  [80000-800ff]  PTE[000-0ff] -----WP 00000-000ff
  [80100-80101]  PTE[100-101] -----P 00100-00101
  [80102-80102]  PTE[102]    -----A----P 00102
  [80103-80105]  PTE[103-105] -----P 00103-00105
  [80106-80106]  PTE[106]    -----A----P 00106
  [80107-80107]  PTE[107]    -----P 00107
  [80108-8010a]  PTE[108-10a] -----WP 00108-0010a
  [8010b-8010b]  PTE[10b]    -----A----WP 0010b
  [8010c-803ff]  PTE[10c-3ff] -----WP 0010c-003ff
[80400-8dffff]  PDE[201-237] -----UWP
  [80400-8dffff]  PTE[000-3ff] -----WP 00400-0dffff
[fe000-ffffff]  PDE[3f8-3ff] -----UWP
  [fe000-ffffff]  PTE[000-3ff] -----WP fe000-ffffff
(qemu)
```

main calls kvmalloc() to create and switch to a page table with the mappings above KERNBASE required for the kernel to run. After paging is enabled, kernel will start executing from virtual address space.

- 3 PDEs are present in Virtual address range 80000-803ff, 80400-8dffff and fe000-fffff.
- All PDEs are User, Writable and Present (as suggested by flags).
- First PDE is accessed. Other two PDEs are not accessed.
- 80000-803ff is level 1 PDE, which contains 1024 PTEs.
- 80400-8dffff contains PDEs in range 201-237 which map physical memory upto PHYSSTOP.
- fe000-fffff contains PDEs that map page tables for I/O devices.

